

Kubernetes



Nicole Heinemann
heinemann@informatik.uni-leipzig.de

What is Kubernetes?

Easy!

“Kubernetes (commonly stylized as K8s) is an open-source container orchestration system for automating software deployment, scaling, and management.”

What is Kubernetes?

What is Kubernetes?

✧ magic ✧

Cloud Native

Docker on Steroids

Very easy and intuitive

PaaS

The thing my SRE friend won't shut up about

Way too complicated

Computing as LEGO

A data center OS

Nightmare fuel

An agnostic platform

The best thing since sliced bread

Declarative Deployment




SwiftOnSecurity
@SwiftOnSecurity

One time I tried to explain Kubernetes to someone.

Then we both didn't understand it.

KÜBERNETES



Containers, that'll fix it.



Kubernetes for beginners.


What could go wrong?

O RLY?

DevOops

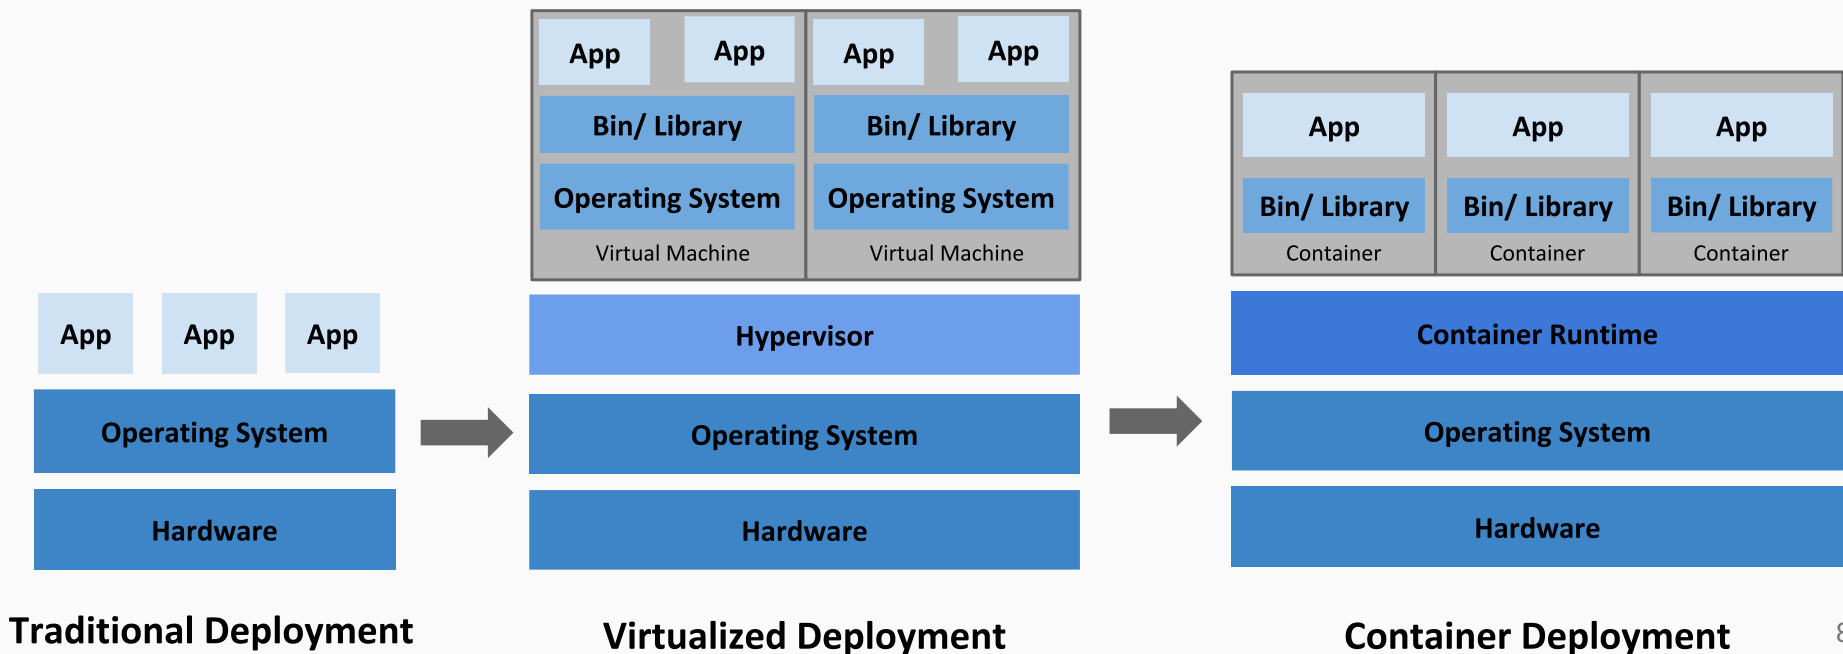


That doesn't seem helpful...




To understand Kubernetes, we first need to understand why it even exists.

A short history of deployment



The old days

- Software running directly on servers
- Maintaining whole infrastructure
- “Throw across the wall deployment”
- Lack of automated deployment
- Conflicting software, no isolation
- Language specific tooling required





Traditional Deployment

The Cloud



VMs in the Cloud



- Infrastructure as a Service
- Easier Scaling, Resilience and Failover
- Dependency Problems still present
- Inefficient Resource Utilization
- "AWS bill problem"
- Still requires
 - Network admin knowledge
 - Storage/DB admin knowledge
 - Linux admin knowledge
 - Cloud specific knowledge



Virtualized Deployment

Containerization

- Solving the “Works on my Machine”
- Integrates nicely in CI/CD
- Less overhead
- Docker run and all is good (?)
- But:
Now deployment requires docker specific skills in addition to all the previous skills



Container Deployment



Help, my docker
container isn't
reachable

```
sudo sysctl -w  
section.tunable=9000  
mount -fvaxuaaa /run/*  
../walk/  
docker pull -f --force  
FORCE=TRUE.'
```

OK LISTEN
CAREFULLY



DID YOU
HIT ENTER?



Yeah I
hit enter

WRONG, I
WASN'T DONE!



```
cat "In the realm of the  
invisible sun, shadows dance with  
the light they cannot touch." >  
/etc/docker/seccomp.conf  
sudo systemctl restart docker
```

Two hard problems looking for a solution...

- Deploying and Maintaining Software Systems
- Managing the underlying Infrastructure

Deploying and Maintaining Software Systems

Containers are great but...

- Scaling
 - Failover
 - Networking
 - Observability
 - Service Discovery
 - Resource Management
 - Logging
 - Access Management
 - ...
- Are still hard!

Managing the underlying Infrastructure

Actual physical computers are...

- A. Weird and non-standard
- B. Bad at computing
- C. Complicated
- D. Rarely in a coherent state
- E. Buggy and prone to failure
- F. All of the above




How does Kubernetes solve this?

- Kubernetes abstracts infrastructure and deployment through heavy use of containerization, virtualized networks and ephemeral storage.
- Deployments are declarative, not imperative
- Users don't need to worry how the definitions are fulfilled.

How does K8s solve this?

The control plane manages a cluster of worker nodes.

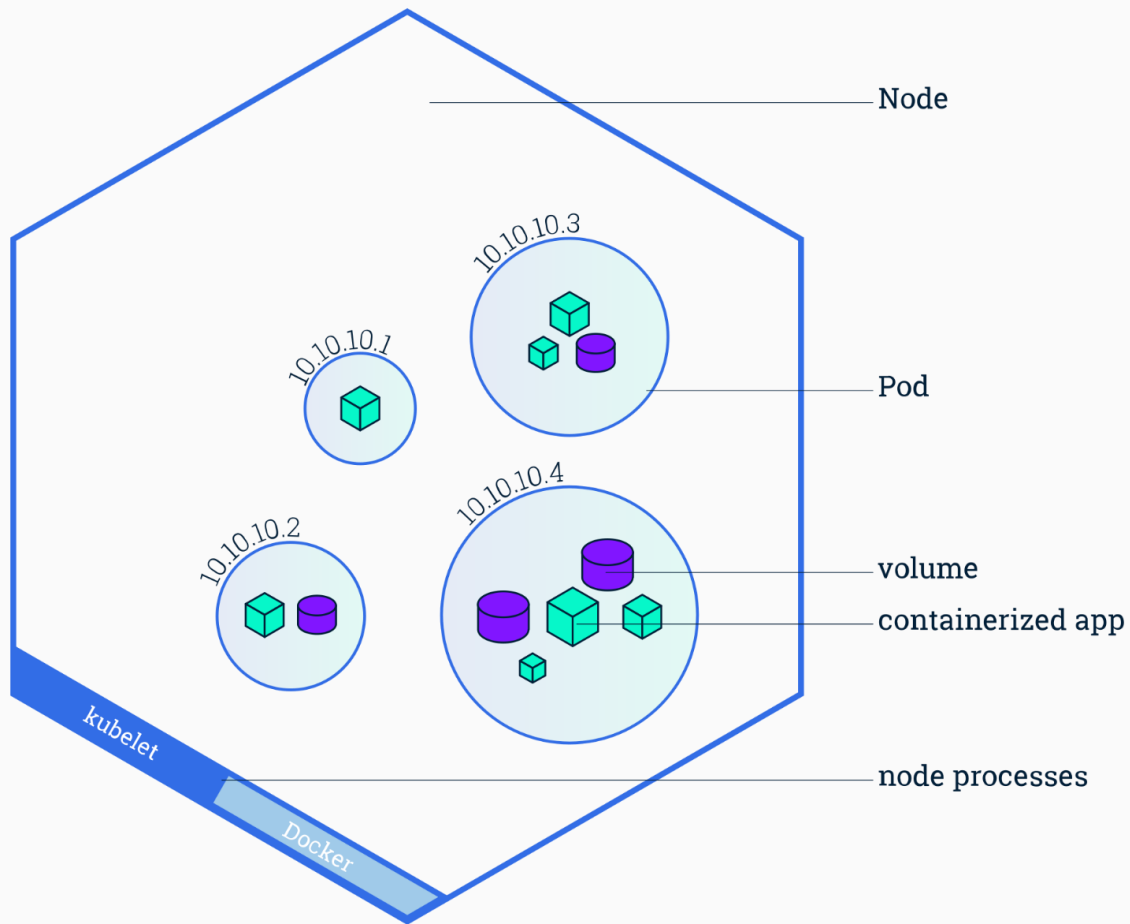
The cluster is exposed as a *virtual platform* to provide resources for the architecture of a Software System.



How does K8s solve this?


Kubernetes defines a set of building blocks (resources).

The characteristics of concrete resources are defined in YAML.



How does K8s solve this?


Higher level resources are built from these primitive resources.



How does K8s solve this?

To deploy a Software System:

1. Write a YAML definition of your architecture
2. **Apply** this definition to a Kubernetes cluster with *kubectl* or the *API*.
3. Kubernetes ensures that those definitions are always fulfilled.



Where do I get one of these clusters?

- Host yourself
- Rent from a cloud vendor



Google Cloud

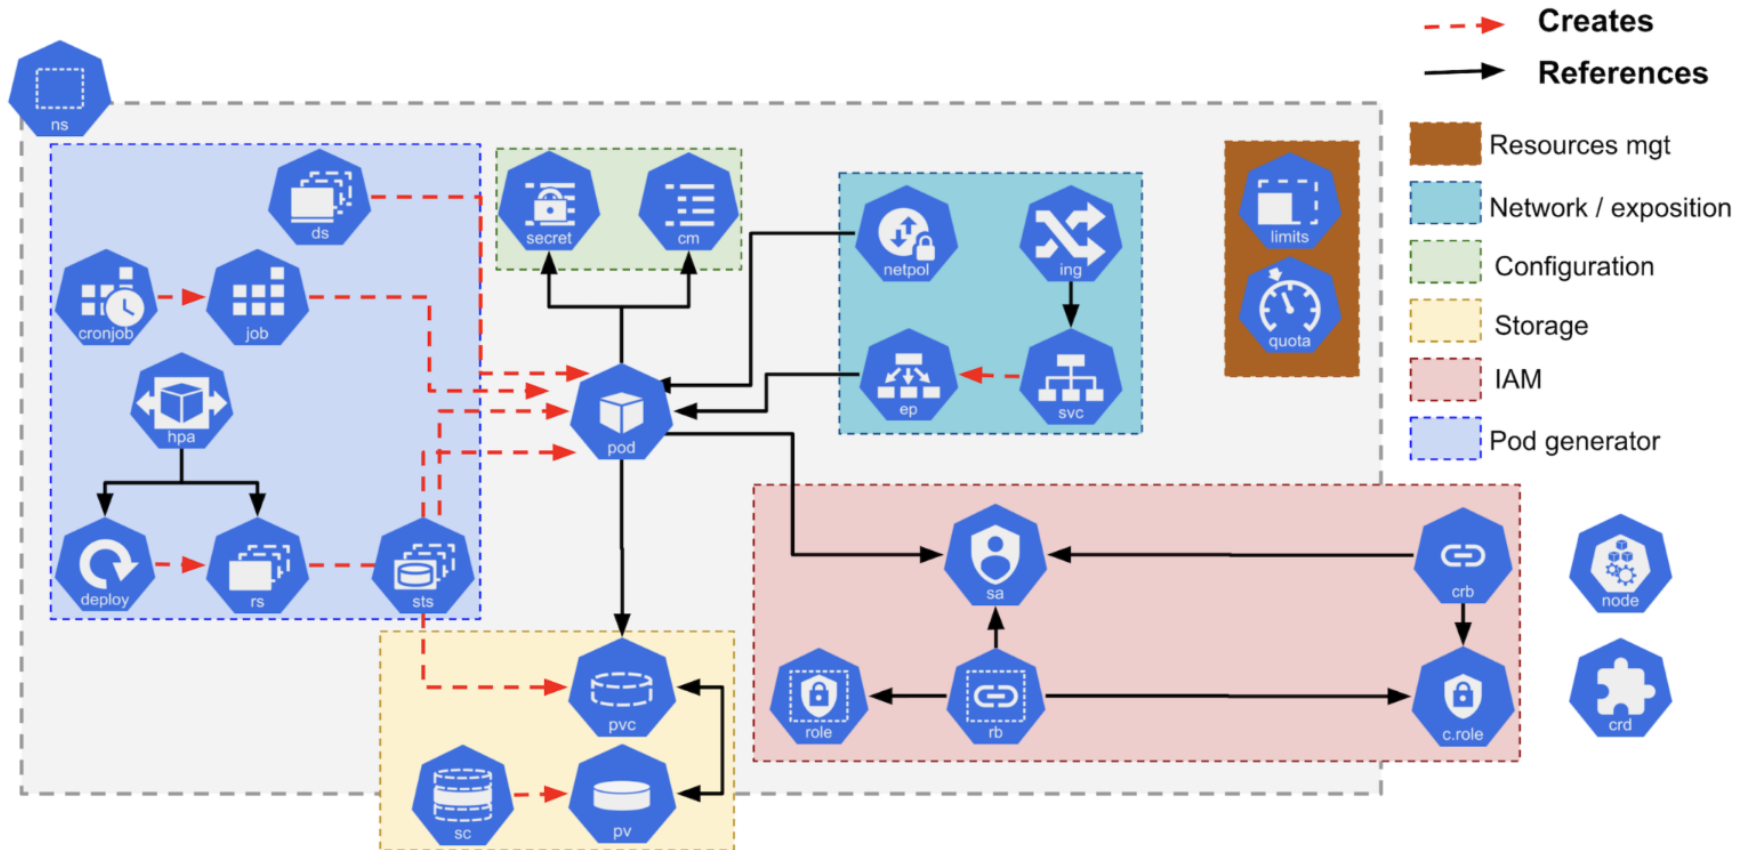


Azure



Kubernetes Resources

Kubernetes Resources Map



Namespace

- Useful for isolating
- Usually one deployment per namespace
- Almost every other resource is attached to a namespace

```
apiVersion: v1
kind: Namespace
metadata:
  name: myspace
```

ConfigMap

- Small (1MB max) immutable configuration storage
- Config itself in a YAML
- Can be mounted as a **Volume**
- Can be accessed through environment variables

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: game-demo
data:
  # property-like keys; each key maps to a value
  player_initial_lives: "3"
  ui_properties_file_name: "user-
interface.properties"
  # file-like keys
  game.properties: |
    enemy.types=aliens,monsters
    player.maximum-lives=5
  user-interface.properties: |
    color.good=purple
    color.bad=yellow
    allow.textmode=true
```

Secret

- Similar to a ***ConfigMap***
- Values are protected
- A change causes a container to be recreated
- Types for common use cases

```
apiVersion: v1
kind: Secret
metadata:
  name: secret-sa-sample
  annotations:
    kubernetes.io/service-account.name: "sa-name"
type: kubernetes.io/service-account-token
data:
  mySecretValue: YmFyCg==
```

Persistent Volume

- Most simple unit of persistent storage
- Requires a storage provider
- Can be attached to ***Pods***
- Only persistent volumes need to be defined!
- Requires a Persistent Volume and a Claim to attach it to a ***Pod***

Ephemeral Volumes are preferred and inferred from *Pod* definitions!

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs
spec:
  capacity:
    storage: 1Mi
  accessModes:
    - ReadWriteMany
  nfs:
    server: nfs.srv.org
    path: "/"
  mountOptions:
    - nfsvers=4.2
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: nfs
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: ""
  resources:
    requests:
      storage: 1Mi
  volumeName: nfs
```

Pod

- Most basic unit of computation
- Contains container(s)
- May contain an *InitContainer*
- References the *Volumes* attached to a container
- Has its own IP address

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pod
spec:
  containers:
  - image: nginx:1.14.2
    ports:
    - containerPort: 80
    volumeMounts:
    - mountPath: /cache
      name: cache-volume
  volumes:
  - name: cache-volume
    emptyDir: {}
```

Deployment

- Consists of metadata and a template for one or more **Pods**
- Defines the number of replicas of said **Pod**
- Often labeled to allow the connection to a **Service**

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2 # tells deployment to run 2 pods
              matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

Service

- Connects a computation resource with a FQDN.
- Every computation resource needs a **Service** to be routable.
- Provides automatically load balancing to the **Pods** in a **Deployments**
- Contains the ports of used to access a **Service**

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

Ingress

- Exposes a **Service** to the external network
- Maps external URLs and FQDNs to internal **Services**
- Usually terminates TLS
- Requires an **Ingress Controller** to be usable.

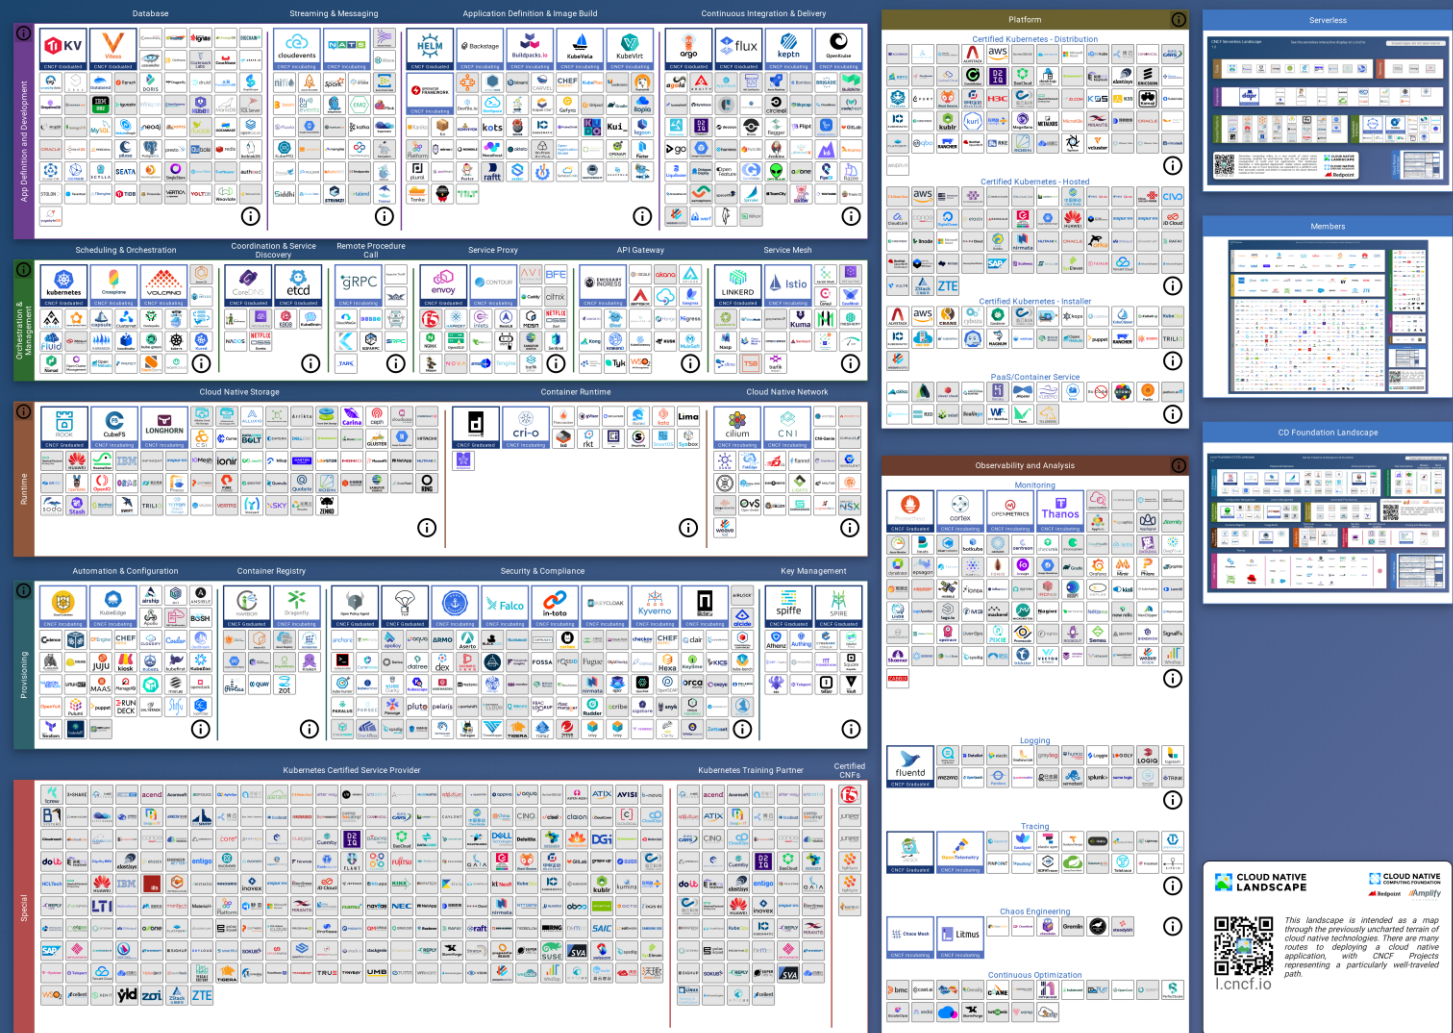
```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: minimal-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx-example
  rules:
    - http:
        paths:
          - path: /testpath
            pathType: Prefix
            backend:
              service:
                name: test
                port:
                  number: 80
```


Custom Resources and Operators

- The Kubernetes is extensible through custom resource definitions (CRD).
- CRDs are managed through operators.
- Operators
 - Plugins for Kubernetes
 - Run on the Cluster with extra privileges
 - Hook into the Kubernetes API

```
apiVersion: acid.zalan.do/v1
kind: postgresql
metadata:
  name: acid-minimal-cluster
  namespace: default
spec:
  teamId: "acid"
  volume:
    size: 1Gi
  numberOfInstances: 2
  users:
    zalando: # database owner
    - superuser
    - createdb
    foo_user: [] # role for application foo
  databases:
    foo: zalando # dbname: owner
  preparedDatabases:
    bar: {}
  postgresql:
    version: "14"
```

The Kubernetes Ecosystem



Helm

- Package Manager for Kubernetes
- Templating YAML Files
- De-Facto Standard to install external software on a cluster



- Provides Certificates
- Letsencrypt Support
- Useful to secure Ingresses
- Prerequisite for many other extensions



Cluster Management

- Provided either by Cloud Vendor or Selfhosted
- Rancher de-facto Standard outside the RedHat Openshift Ecosystem.



Networking

- Provided via CNI-Plugins
- Essential components in Kubernetes
- Enable network connectivity for containers within a cluster
- Manage IP addresses, routes, etc.
- Facilitate communication between containers, pods, and external networks



cilium



PROJECT
CALICO

Storage

- Required for Persistent Volumes
- Require a CSI Driver
- Externally Provided by
 - Cloud Vendor
 - NFS, iSCSI, etc.
- Hosted in the Cluster itself
 - Longhorn
 - Rook (Ceph)



Ingress Providers

- Required for Ingresses
- Can bring their own extensions to Kubernetes networking
- Less relevant in the cloud as an external Load Balancer is often provided



Continuous Deployment

- Specifically designed for Kubernetes
- Almost always following a GitOps approach



Service Mesh

Provide advanced networking capabilities:

- Traffic isolation
- mTLS
- Traffic Observability
- Rate Limiting
- A/B test routing
- Additional observability



Istio



LINKERD

Observability

- Collecting metrics:
 - Resource utilization
 - Response time
 - Outage
- Sending alerts



Prometheus

Logging

- Collecting logs from all Containers
- All Information accessible in one location
- Popular solutions:
 - ELK stack
 - EFK stack



Container Registry

- "Dockerhub at home"
- Faster image pulls and caching
- Provide additional capabilities:
 - Signing
 - Vulnerability scans
 - Autocleaning



HARBOR



Dragonfly

Kubernetes for AI enable Software

Why Kubernetes for AI enabled Software?

In deploying AI enabled software systems, the issues intrinsic to traditional systems are frequently magnified:

- The need for scalability to accommodate larger models
- Increased frequency of deployments due to regular updates
- A higher rate of rollbacks due to potential issues with model quality
- The expense of inference necessitates efficient hardware utilization

Kubernetes design philosophy



Declarative over
imperative
deployments



Treat deployments
like cattle, not pets



Operations must be
atomic



Resilience and
redundancy aren't
optional



Extensibility over
batteries included

GitOps and Kubernetes

What is GitOps?

Builds on

- Monorepo
- Infrastructure/Configuration as Code
- CI/CD
- Deployment as Declaration

To enable a **Single Source of Truth** for Deployment




Why GitOps

Single Source of Truth enables:

- Unified Development Flow
- Coherent State
- Easier Rollbacks and Debugging
- Centralized resource management

How does GitOps work?



Challenges of GitOps

- Initial Learning Curve
- Secrets Management
 - Secrets Provider (Hashicorp Vault, etc.)
- Complex dynamic deployments
 - Dedicated deployment pipelines
- Requires Strict Git Discipline



So, how do you work with Git?



Kubernetes as a Platform for AI

Why extend Kubernetes

- Flexible on-demand scaling
- Isolation through Containerization
- Workload management

→ Needed for managing big data, training computationally intensive models and scalable inference.

- Unified Developer Experience
- Extendible API
- Operator Framework

→ Makes the life of developers and users much easier

Kubernetes as an “AI Platform”-Platform

Many AI Development and Deployment Tools use Kubernetes at their core or are adapted to fit into the ecosystem



Apache
Airflow



FEAST



Kubeflow



jupyterhub

mlflowTM



JupyterHub

- Multi-User Notebooks with simultaneous access
- Secured, isolated workspaces
- Allows customizable setups,
supporting different kernels and configurations
- Enables Data Scientists to access large computational resources
(including GPUs) on their end device.



Kubeflow

- Kubernetes-Native
- All in one platform
- End-to-End ML workflow from data ingestion and preprocessing to model training, serving, and monitoring
- Component-Based Architecture:
 - Jupyter notebooks
 - Training jobs operators
 - Serving stack for deploying models



Kubeflow



Extremely hard to set up
outside of cloud vendors



Apache Spark

- Unified computing engine for big data processing and analytics
- Supports SQL, streaming and ML
- In-memory processing
- Bindings for Java, Scala and Python
- Adapted to Kubernetes to utilize its resource management capabilities



Apache Airflow

- Platform to programmatically author, schedule and monitor tasks in pipelines
- Pipelines can be dynamically triggered
- Utilizes directed acyclic graphs (DAGs)
- Extensible via plugins
- Scales task execution across multiple worker



...and many more



Security in Containers and Kubernetes

Security issues of Containers

- Containers are “Root Filesystems in a Tarball”
 - Libraries frozen in time
- Still huge attack surface
- Horizontal escalation
- Supply chain risk
- Leaky isolation

Mitigation: Minimalist containers

Utilizing a multi-stage build approach:

1. Build in a rich environment
2. Execute artifacts in a minimalist container
 - Alpine
 - Distroless



alpine
Linux

Mitigation: Reduced Privileges

- Running as non-root in a container
- Setting CPU and memory limits
- Enforcing a security confinement (Seccomp + SELinux)
- Rootless containers with user privileges

Mitigation: Signing

- Signing containers to provide integrity and authorship
- Can prevent supply chain attacks
...but only if those signatures are enforced!
- Long Term Certificates can cause further problems
→ Sigstore and Cosign



Further Mitigations

- Containers as actual VMs:

- Firecracker
- Kata Containers
- gVisor

- Container scanning?

→ Highly unreliable



Kubernetes Demo